

① Devops - Fundamentals

- ② CI/CD Pipeline.
- ③ Cloud engineering using Infra.
- ④ Docker.
- ⑤ Terraform.

Penelji :-

- ① Kubernetes (K8)
- ② Observability (Splunk, Dynatrace)
- ③ Best Practices & Governance.

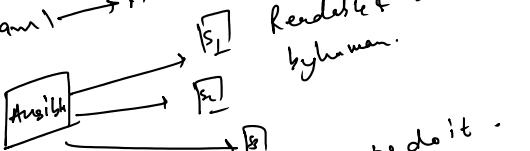
- ① CI/CD Best Practices
- ② Introduction to Ansible for Configuration Management -

① Ansible

- Ansible - F/T Automation engine used for Configuration Management -
- ① F/T Automation engine used for Configuration Management -
 - ② Red Hat → IBM.
 - ③ Core features:-
 - ① Written in Python.
 - ② uses YAML for Config (Playbooks)
 - ③ SSM to communicate with target machine.
 - ④ Push-based model

Why? Does not require agent or daemon to be installed in target Windows.

- ① Agents! - Does not require agent or daemon to be installed in target Windows.
- ② Idempotent, running playbook → same result every time.
- ③ Simple yaml Syntax - yaml → Automate the task → playbook



What to Achieve & Not how to do it.

Declarative style

Provisioning, Config. mang., App. Deployment

Orchestration, CD, IaC

chef, puppet, SaltStack

↓

Ruby Ruby

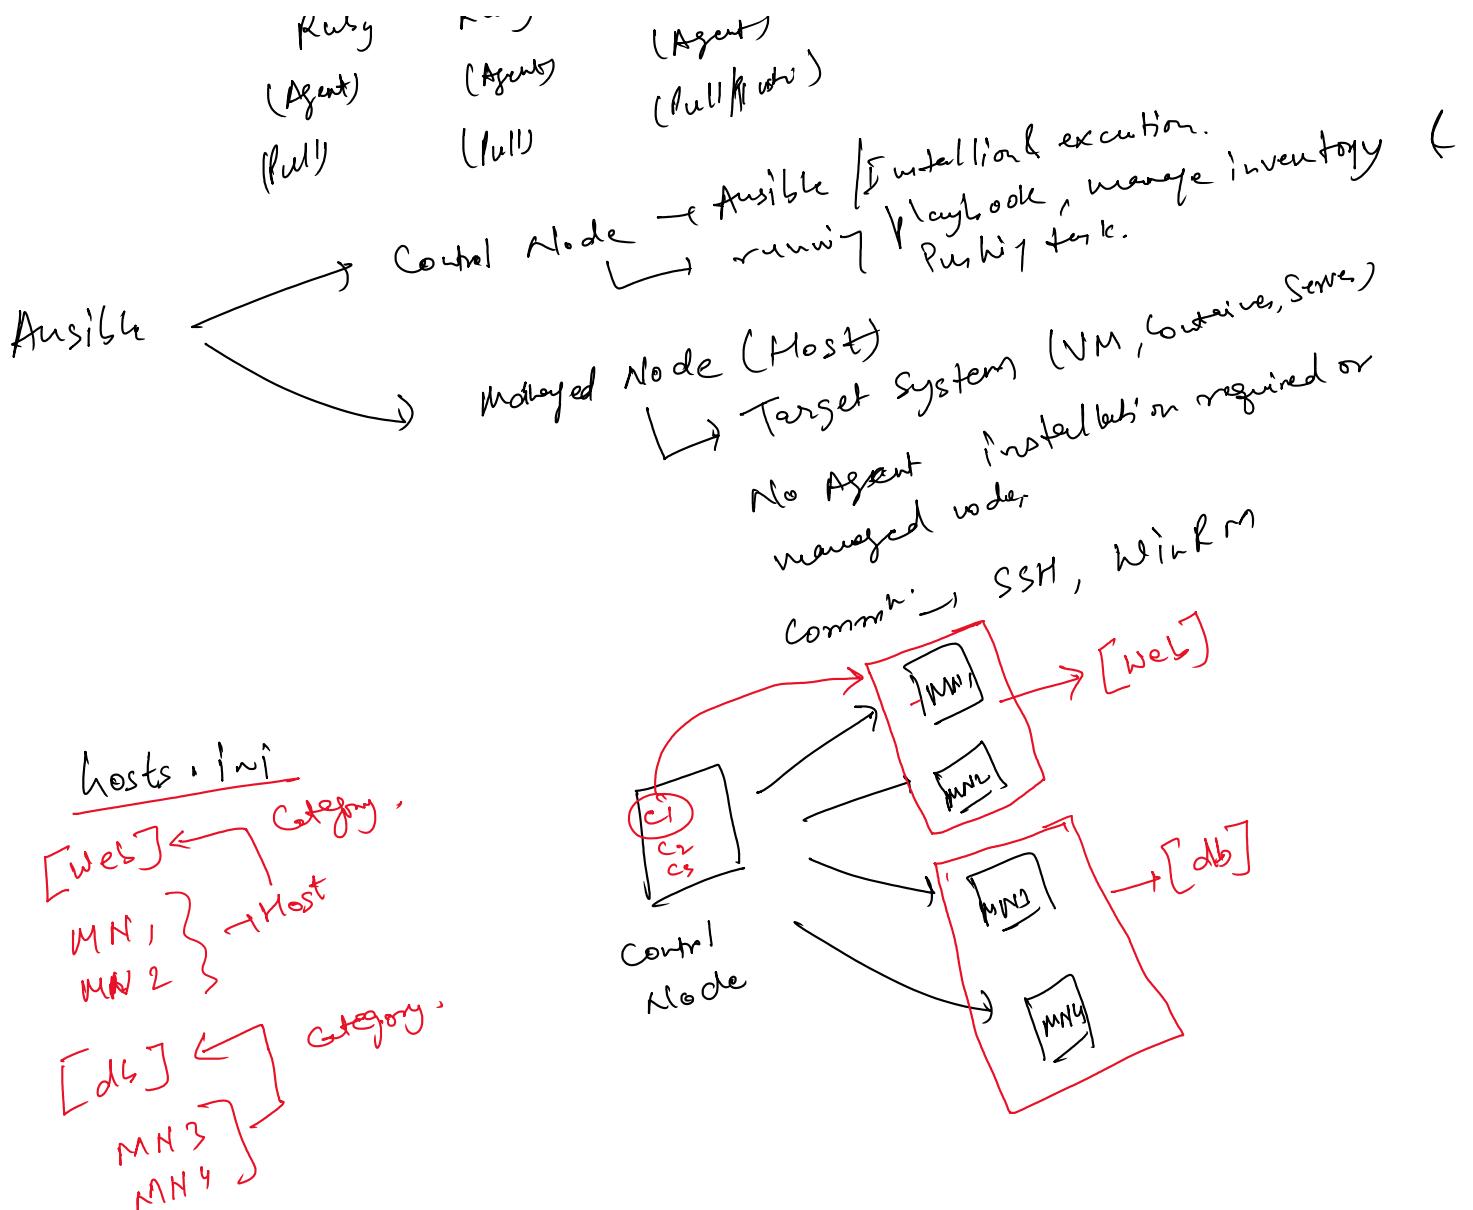
(Agent)

(Agent)

YAML & Python

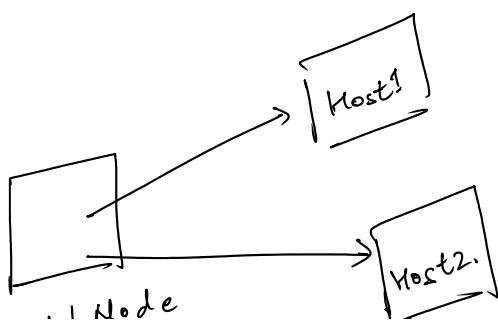
(Agent)

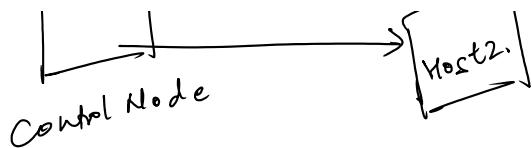
(null/inf)



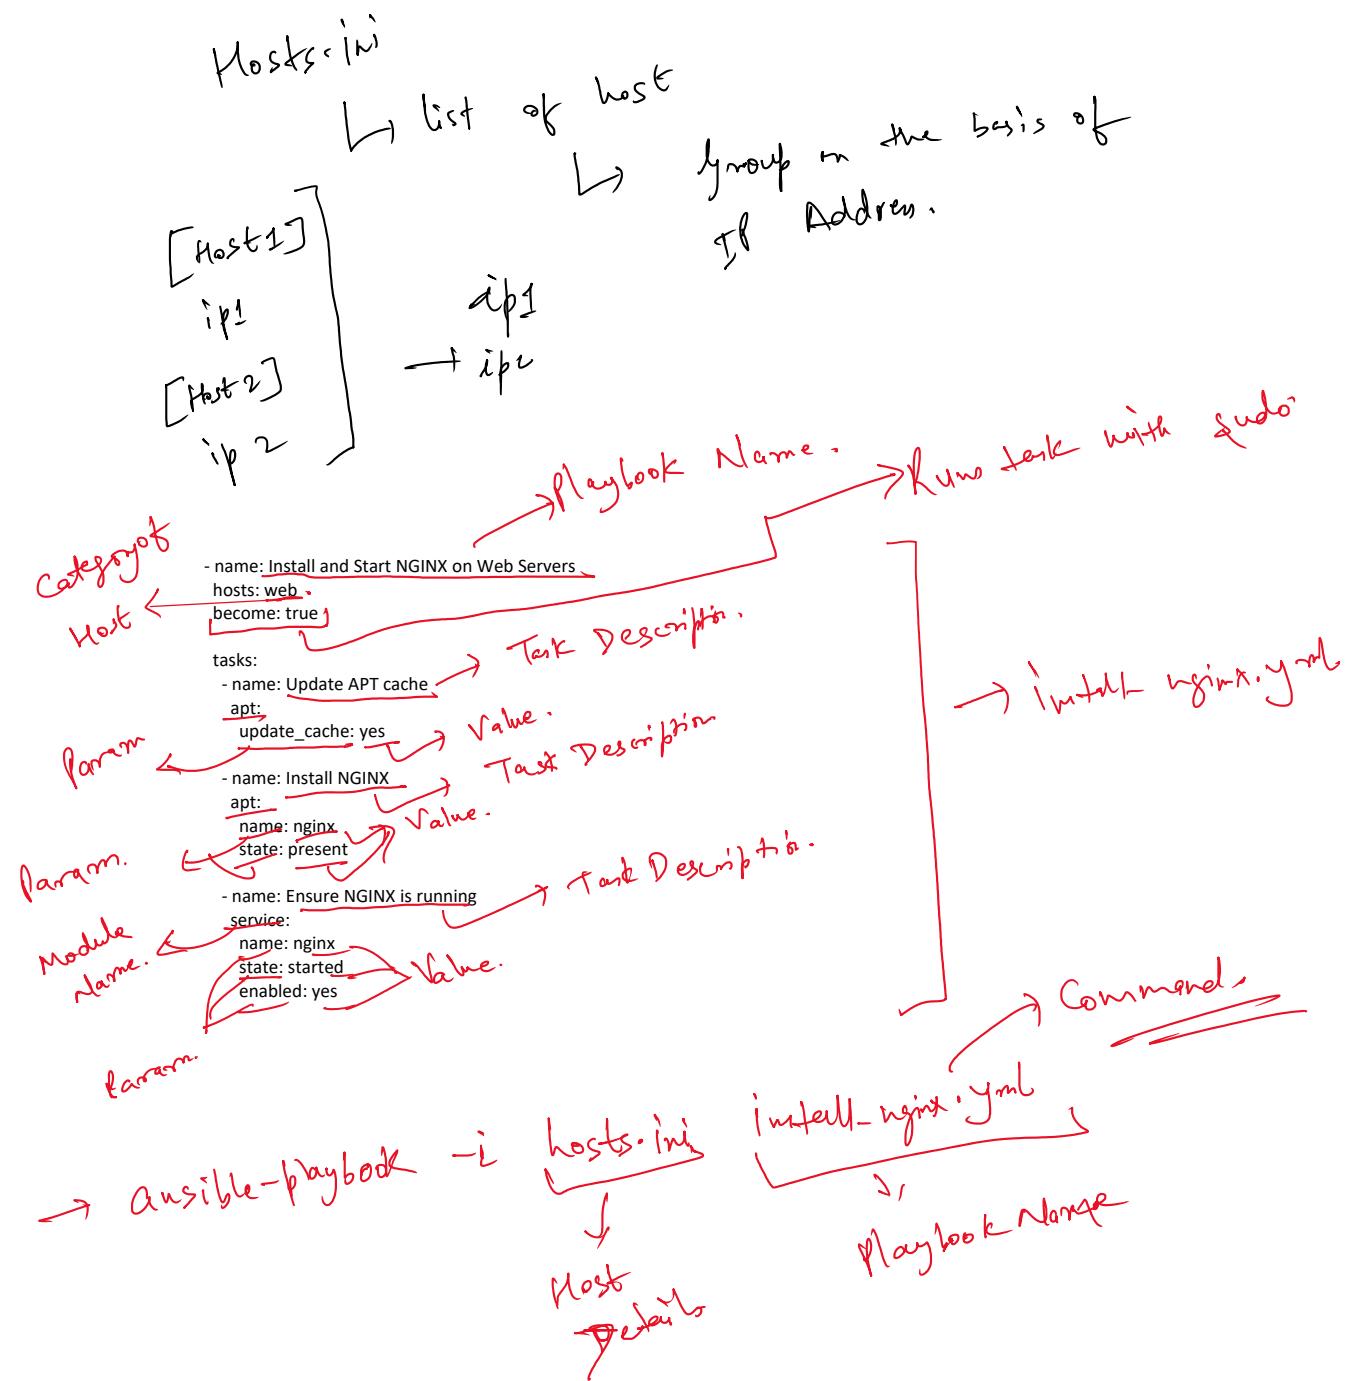
Module → Unit of work used in Ansible.

- ① Core Module → Ansible Team
 - ② Community Module → via Galaxy.
 - ③ Custom module → User-defined Python Script
 - ④ Action Plugins → Specialized Module, wrapper for Playbook.



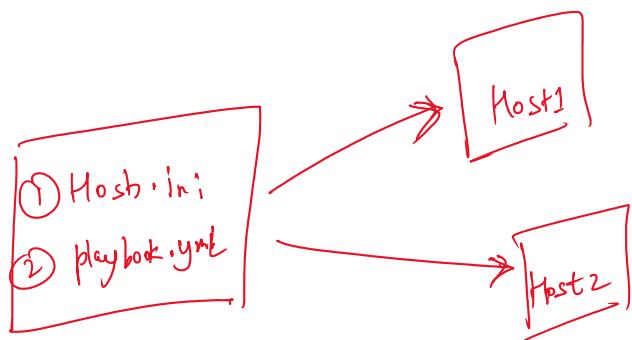


```
- name: Play Name
hosts: group_name_or_ip
become: true      # Run with sudo
vars:
  var1: value1
tasks:
  - name: Task 1 description
    module_name: -
      param1: value1
      param2: value2
  - name: Task 2 description
    module_name:
      param: value
```

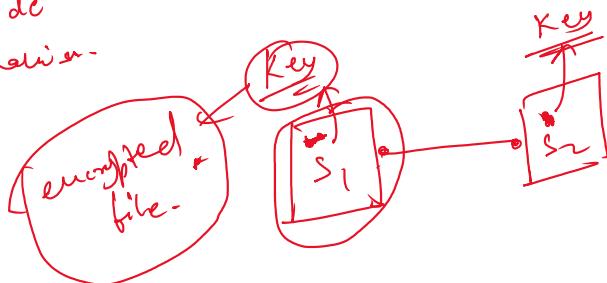


apt, yum → install package
 copy, template → Manage files
 Service → Manage Service
 user, group → Manage users

file → Set permission / ownership
 shell, command → Shell / command task
 debug → Print info during execution



Control Mode
 Pull Mechanism



* Pipeline Governance & Approval in CI/CD:

Pipeline governance → Controls, policies & review mechanism.

- ① Compliant with org. rules.
- ② Secured & reviewed.
- ③ Audit & traceable.
- ④ Approved by stakeholders.

Why?

① Security - prevent unauthorized or unsafe code
I... reaching production.

I... req. like HIPAA, PCI-DSS

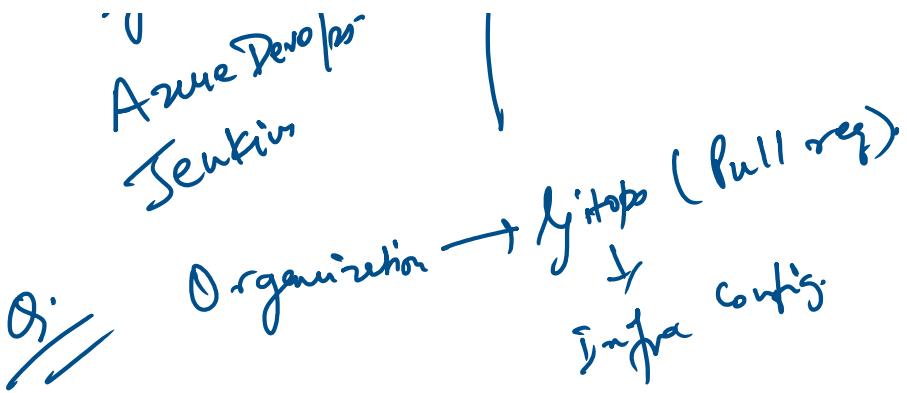
- ① Security - prevent for reusing product...
Satisfy regulatory req. like HIPAA, PCI
- ② Compliance - Maintain logs
- ③ Auditability - Reduce failure | Breach.
- ④ Risk Management - Enforce ownership & transparency for code changes.
- ⑤ Team Accountability - Transfer ownership for code changes.

Component :-

- ① RBAC -
Manual Approval | Approval later.
- ② Policy as Code -
Change management Integrations
- ③ Audit log & Signature
like Signoff, Cosign.
- ④
- ⑤

github Action
github CI/CD
Azure DevOps

Argo CD
Spinnaker



Question: An organization employs GitOps with Kustomize overlays to manage its infrastructure configurations across development, staging, and production environments. A critical security update needs to be propagated immediately to all environments, but the GitOps pipeline mandates pull request reviews for every environment branch. Which GitOps strategy best accelerates this update while maintaining auditability and preventing configuration drift?

Options:

- A. Commit the update directly to the production branch first, then backport the change to development and staging branches afterward
 - instability
 - Config. drift
 - + very risky
- B. Use git cherry-pick to selectively apply the update commit across all environment branches promptly
 - (dev, stage or prod.)
 - ↳ efficient
- C. Create a hotfix branch for the update, merge it immediately into production, and then sequentially merge it into other environment branches
 - Slow,
- D. Deploy the update to staging first, perform thorough testing, and then promote the tested artifact to development and production environments
 - Too slow,

Question Breakdown

- ① Gitops with Kustomize overlay
- ② dev, stage, prod - Separate git branch.
- ③ Critical Security update → immediately to all the environment.
- ④ Pull req. (required)

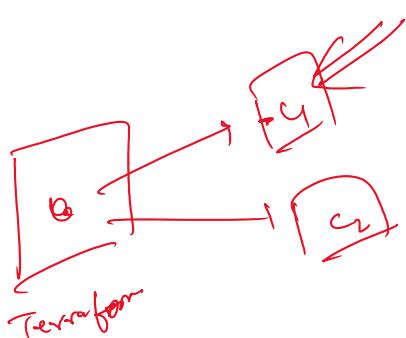
An organization employs Infrastructure as Code (IaC) to manage a complex cloud environment. During a recent audit, discrepancies were found between the desired state defined in the IaC configuration and the actual infrastructure state, indicating drift. To effectively address this issue with minimal manual intervention while ensuring configuration consistency, the team needs to implement a reconciliation strategy that automatically detects and corrects these deviations. What is the most appropriate strategy to resolve this drift?

Options:

→ Reconcile

Options:

- A. Implement scheduled reconciliation that runs periodically to detect and automatically correct configuration deviations, maintaining alignment efficiently
- B. Use imperative scripts triggered manually to identify and correct configuration discrepancies as necessary, relying on human initiation
↳ *No Automation*.
- C. Monitor configuration drift manually and apply all fixes through formal change requests without automation
↳ *Without Automation*.
→ *missed delay, missed issue, audit failure*
- D. Employ continuous reconciliation that perpetually monitors the infrastructure and instantly corrects deviations to ensure immediate compliance
↳ *Complex, Uncommon*



↳ *fac*

- ① Code Reuse
② code duplication reduction -
③ easier control
④ faster update.

→ *Not recommended*

An organization is setting up a cloud infrastructure using Terraform involving multiple services such as databases and networks. The team aims to improve code reusability for better organization and maintainability. Which approach best enhances code reuse and reduces code duplication, enabling easier control and faster updates?

Options:

- A. Use a single, monolithic Terraform configuration file that manages all infrastructure components in one place
- B. Duplicate similar code blocks across various Terraform configuration files, which risks inconsistencies during future updates
- C. Implement Terraform modules to encapsulate and reuse common infrastructure components effectively across configurations
↳ *Reuse, consistent, efficient*
- D. Apply a complex naming convention to resources instead of adopting proper modularization techniques X